

AI Resume Screening & Ranking — Results

Generated 2026-09-23 17:15 | LLM used: True | duration: 10.35s

Batch summary

Total resumes	50	Eligible	41
Parsed OK	50	Rejected	9
Parse failed	0	Duplicates skipped	0
LLM scored (survivors)	41	Skipped by hard gate	9
GitHub enriched	6	GitHub no profile	25
GitHub tagged (rate-limited)	10	GitHub tagged (error)	0

Shortlist — eligible candidates (41)

Det = deterministic baseline (no LLM) • **Final** = with LLM project scoring • Δ = LLM contribution • **AI** shows det→llm depth. Every other category is deterministic and unchanged.

#	Candidate	Det	Final	Δ	AI (det→llm)	PY	Cloud	GH	Pen
1	Yash Maini	74.2	99.5	+25.3	15→40	30	15	10	0
2	PRAJWAL A S	84.7	88.2	+3.5	36→39	30	6	8	0
3	V Sree Raghu Vardhan	60.0	88.0	+28.0	10→38	30	15	0	0
4	Hari Shanker Sharma	74.6	86.5	+11.9	25→36	30	15	0	0
5	ARUNIMA SAHA	59.9	85.6	+25.7	10→36	30	15	0	0
6	ABHINAV MISHRA	76.5	85.0	+8.5	22→30	27	13	10	0
7	Gourab Das	62.8	85.0	+22.2	8→30	30	15	5	0
8	Vivek Chimnani	77.6	83.3	+5.7	23→29	29	15	5	0
9	Prathamesh Patil	83.7	82.2	-1.5	34→35	30	15	0	-3
10	Sumaiya Sultana Shaik	65.2	81.7	+16.5	28→32	30	15	0	0
11	Thejeshwaar Paasila	73.2	81.5	+8.3	18→27	30	15	5	0
12	Kiran Naregal	61.5	78.8	+17.3	16→36	24	13	5	-3
13	Nevil Jasmin Mehta	53.1	78.6	+25.5	16→29	30	15	0	0
14	Aditi Kala	57.1	78.0	+20.9	6→33	29	15	5	-6
15	Arjun Kumar	73.8	75.2	+1.4	25→27	30	8	10	0
16	RAJ THUMAR	70.7	75.0	+4.3	25→29	30	11	0	0
17	Vaibhav Wakde	61.6	74.1	+12.5	20→32	30	7	0	0
18	SHUBHANSHU NAVADIYA	53.4	70.2	+16.8	17→36	30	4	0	-3
19	PAVANI A	61.9	69.9	+8.0	16→21	30	13	5	0
20	Jyandeep Baishya	62.7	69.7	+7.0	28→35	25	5	0	0
21	ANNISHA S	46.1	69.1	+23.0	16→26	30	7	5	0
22	Agam Jain	70.1	68.7	-1.4	27→25	30	10	0	0
23	Vivek Shekhawat	51.1	68.0	+16.8	13→29	24	5	5	0
24	PRIYANSHU KUMAR	53.8	65.8	+12.0	25→37	23	3	0	0
25	HARSHA A	55.1	64.3	+9.1	17→26	30	7	0	0

... and 16 more eligible candidate(s).

How the final score is built (deterministic vs LLM)

Every category is computed deterministically from resume evidence. The LLM refines **only** the AI / Agentic project-depth category, scoring each AI project 0-10 with a rationale and cited evidence. The **final** score is the deterministic score with the AI-depth category replaced by the LLM judgement; all other categories are identical. For the score breakdown below, the AI-depth row shows what the deterministic signal estimate gave versus what the LLM decided.

#1 Yash Maini — final 99.5 (deterministic 74.2, LLM +25.3)

Category	Deterministic	With LLM	Notes
AI project depth (40)	14.7	40.0	CopyGuard: depth 5.0/10 — The project demonstrates solid production-grade infrastructure and backend engineeri
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	9.5	9.5	last public activity ~3d ago; 2 event(s) in 90d; relevant repos: face-
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	74.2	99.5	LLM contribution +25.3

Project: CopyGuard: A serverless GenAI code detection platform built on AWS Lambda and Bedrock. It provides low-latency detection for burst traffic with full infrastructure-as-code implementation.

#2 PRAJWAL A S — final 88.2 (deterministic 84.7, LLM +3.5)

Category	Deterministic	With LLM	Notes
AI project depth (40)	35.7	39.2	AI Voice Appointment Assistant: depth 7.0/10 — The project demonstrates strong backend integration and tool-ca
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	6.0	6.0	deterministic (unchanged)
GitHub (10)	8.0	8.0	last public activity ~47d ago; relevant repos: voice-ai-backend, pytho
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	84.7	88.2	LLM contribution +3.5

Project: AI Voice Appointment Assistant: A real-time, multilingual voice AI agent for healthcare appointment management. It utilizes WebRTC for low-latency interaction and integrates with a PostgreSQL database via FastAPI to ensure grounde

#3 V Sree Raghu Vardhan — final 88.0 (deterministic 60.0, LLM +28.0)

Category	Deterministic	With LLM	Notes
AI project depth (40)	10.0	38.0	AI-driven enterprise HRMS workflows: depth 10.0/10 — This is a highly sophisticated production-grade system fe
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	0.0	0.0	No GitHub profile on resume
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	60.0	88.0	LLM contribution +28.0

Project: AI-driven enterprise HRMS workflows: A complex system of multi-step AI agents integrated with enterprise APIs, semantic retrieval, and document processing to automate HR tasks. It features structured output validation, Redis-backe

#4 Hari Shanker Sharma — final 86.5 (deterministic 74.6, LLM +11.9)

Category	Deterministic	With LLM	Notes
AI project depth (40)	24.6	36.5	500Agents: depth 9.0/10 — Demonstrates high-level agentic orchestration using LangGraph, multi-provider routin
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	0.0	0.0	No GitHub profile on resume
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	74.6	86.5	LLM contribution +11.9

Project: Sasta Signals: A distributed price monitoring system that tracks products across 7 major Indian e-commerce platforms using automated web scraping. It features real-time price drop detection with notifications and a robust backend

#5 ARUNIMA SAHA — final 85.6 (deterministic 59.9, LLM +25.7)

Category	Deterministic	With LLM	Notes
AI project depth (40)	10.0	35.7	Multi-Agent Travel Planning System: depth 9.0/10 — This is a sophisticated multi-agent system featuring comple
Python & backend (30)	29.9	29.9	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	0.0	0.0	No GitHub profile on resume
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	59.9	85.6	LLM contribution +25.7

Project: Multi-Agent Travel Planning System: An agentic travel planning and booking system that uses LangGraph to orchestrate specialized nodes for flight/hotel search and booking. It features human-in-the-loop workflows, persistent state

#6 ABHINAV MISHRA — final 85.0 (deterministic 76.5, LLM +8.5)

Category	Deterministic	With LLM	Notes
AI project depth (40)	21.6	30.1	Construct IDE: depth 8.0/10 — The project demonstrates sophisticated agentic orchestration and state managemen
Python & backend (30)	26.6	26.6	deterministic (unchanged)
Cloud / full-stack (15)	13.2	13.2	deterministic (unchanged)
GitHub (10)	10.0	10.0	last public activity ~0d ago; 60 event(s) in 90d; relevant repos: hack
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	76.5	85.0	LLM contribution +8.5

Project: Construct IDE: An AI-native teaching IDE that generates personalized coding lessons, implementation tasks, and concept checks. It maintains persistent learner state and session continuity across project steps using LangGraph.

#7 Gourab Das — final 85.0 (deterministic 62.8, LLM +22.2)

Category	Deterministic	With LLM	Notes
AI project depth (40)	8.0	30.1	Tool-augmented AI agents: depth 8.0/10 — This project utilizes advanced orchestration frameworks and tool-call
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	4.8	4.8	last public activity ~2d ago; 1 event(s) in 90d
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	62.8	85.0	LLM contribution +22.2

Project: SMILE: A full-stack platform built with Django and React to manage donor and patient data. It includes automated infrastructure deployment using Docker and CI/CD pipelines.

#8 Vivek Chimnani — final 83.3 (deterministic 77.6, LLM +5.7)

Category	Deterministic	With LLM	Notes
AI project depth (40)	23.2	28.9	AI Research Copilot (LangGraph): depth 8.0/10 — The project demonstrates advanced agentic design through multi
Python & backend (30)	29.4	29.4	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	5.0	5.0	GitHub rate limit already hit this run
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	0.0	thin-wrapper / tutorial
TOTAL	77.6	83.3	LLM contribution +5.7

Project: AI Research Copilot (LangGraph): A research assistant utilizing a multi-node workflow for planning, research, and report generation. It features session persistence, conditional routing, and modular backend APIs.

#9 Prathamesh Patil — final 82.2 (deterministic 83.7, LLM -1.5)

Category	Deterministic	With LLM	Notes
AI project depth (40)	33.7	35.2	Adaptive Agentic RAG - Multi-Agent Research System: depth 9.0/10 — The project demonstrates high-level enginee
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	0.0	0.0	No GitHub profile on resume
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	0.0	-3.0	thin-wrapper / tutorial
TOTAL	83.7	82.2	LLM contribution -1.5

Project: Adaptive Agentic RAG - Multi-Agent Research System: A multi-agent research system using LangGraph for conditional routing between vector stores and web search. It features self-correction nodes for hallucination detection, persist

#10 Sumaiya Sultana Shaik — final 81.7 (deterministic 65.2, LLM +16.5)

Category	Deterministic	With LLM	Notes
AI project depth (40)	28.1	31.6	Privacy Risk Intelligence Platform: depth 8.0/10 — The project demonstrates a sophisticated multi-layer archit
Python & backend (30)	30.0	30.0	deterministic (unchanged)
Cloud / full-stack (15)	15.0	15.0	deterministic (unchanged)
GitHub (10)	0.0	0.0	No GitHub profile on resume
Engineering depth (5)	5.0	5.0	deterministic (unchanged)
Penalties	-13.0	0.0	thin-wrapper / tutorial
TOTAL	65.2	81.7	LLM contribution +16.5

Project: Privacy Risk Intelligence Platform: A 3-layer privacy-risk analysis system that combines NLP for PII detection, a Random Forest classifier for risk scoring, and a LangChain-orchestrated LLM layer for remediation advice. The system

Rejected candidates (9) — and why

Eligibility is a hard, deterministic filter: a candidate must show **Python** AND **AI/agent** evidence. Reasons name only the unmet requirement — never the presence of an “unwanted” skill. *matched_skills* is the evidence that was found.

Candidate	Rejection reason(s)	Matched skills (what we did find)
AABHAR GUPTA	No evidence of Python stack; No AI/agent project evidence	PostgreSQL, MongoDB, Redis, REST API, Node.js, Express, JavaScript, GCP, CI/CD, React
ANJALI PATIL	No evidence of Python stack; No AI/agent project evidence	PostgreSQL, MySQL, Redis, SQL, REST API, Java, Spring Boot, JavaScript, Docker, Testing
Gayathri Palraj	No AI/agent project evidence	Python, MySQL, MongoDB, SQL, REST API, Node.js, Express, Java, TypeScript, JavaScript
Harshit Kumar Mishra	No evidence of Python stack; No AI/agent project evidence	PostgreSQL, MongoDB, Redis, SQL, GraphQL, WebSockets, Node.js, Express, C++, TypeScript
Kartikay Sinha	No evidence of Python stack; No AI/agent project evidence	MySQL, MongoDB, SQL, REST API, Node.js, Express, C++, TypeScript, JavaScript, Docker
MD AZIZUL ISLAM	No evidence of Python stack	OpenAI API, RAG, Embeddings, pgvector, LLM, PostgreSQL, MongoDB, Redis, SQL, REST API
Mukul Pal	No evidence of Python stack	OpenAI API, LLM, PostgreSQL, MongoDB, Redis, REST API, WebSockets, Node.js, Express, TypeScript
SHIVAM RAJ	No evidence of Python stack; No AI/agent project evidence	REST API, Node.js, Express, TypeScript, JavaScript, Azure, Docker, CI/CD, Vercel, Render
Shreya Chauhan	No AI/agent project evidence	Python, MySQL, MongoDB, Redis, SQL, REST API, Kafka, Node.js, Express, Java

Eligibility is a hard, deterministic filter (Python AND AI/agent evidence). Scores are evidence-source weighted; the LLM contributes AI project depth with cited evidence. GitHub rate-limited/errored profiles are tagged and given a neutral placeholder.