

AI Resume Screening & Ranking — Results

Generated 2026-09-23 16:59 | LLM used: True | duration: 419.37s

Batch summary

Total resumes	50	Eligible	41
Parsed OK	50	Rejected	9
Parse failed	0	Duplicates skipped	0
LLM scored (survivors)	41	Skipped by hard gate	9
GitHub enriched	6	GitHub no profile	25
GitHub tagged (rate-limited)	10	GitHub tagged (error)	0

Shortlist — eligible candidates (41)

#	Candidate	Total	AI	PY	Cloud	GH	Eng	Pen
1	Yash Maini	99.5	40	30	15	10	5	0
2	PRAJWAL A S	88.2	39	30	6	8	5	0
3	V Sree Raghu Vardhan	88.0	38	30	15	0	5	0
4	Hari Shanker Sharma	86.5	36	30	15	0	5	0
5	ARUNIMA SAHA	85.6	36	30	15	0	5	0
6	ABHINAV MISHRA	85.0	30	27	13	10	5	0
7	Gourab Das	85.0	30	30	15	5	5	0
8	Vivek Chimnani	83.3	29	29	15	5	5	0
9	Prathamesh Patil	82.2	35	30	15	0	5	-3
10	Sumaiya Sultana Shaik	81.7	32	30	15	0	5	0
11	Thejeshwaar Paasila	81.5	27	30	15	5	5	0
12	Kiran Naregal	78.8	36	24	13	5	4	-3
13	Nevil Jasmin Mehta	78.6	29	30	15	0	5	0
14	Aditi Kala	78.0	33	29	15	5	2	-6
15	Arjun Kumar	75.2	27	30	8	10	1	0
16	RAJ THUMAR	75.0	29	30	11	0	5	0
17	Vaibhav Wakde	74.1	32	30	7	0	5	0
18	SHUBHANSHU NAVADIYA	70.2	36	30	4	0	2	-3
19	PAVANI A	69.9	21	30	13	5	1	0
20	Jyandeep Baishya	69.7	35	25	5	0	4	0
21	ANNISHA S	69.1	26	30	7	5	1	0
22	Agam Jain	68.7	25	30	10	0	4	0
23	Vivek Shekhawat	68.0	29	24	5	5	5	0
24	PRIYANSHU KUMAR	65.8	37	23	3	0	2	0
25	HARSHA A	64.3	26	30	7	0	1	0

... and 16 more eligible candidate(s).

Score breakdown — top candidates

#1 Yash Maini — 99.5/100

Project: CopyGuard: A serverless GenAI code detection platform built on AWS Lambda and Bedrock. It provides low-latency detection for burst traffic with full infrastructure-as-code implementation.

GitHub: last public activity ~3d ago; 2 event(s) in 90d; relevant repos: face-verification-resnet50, pg-reliability-agent

- CopyGuard: depth 5.0/10 — The project demonstrates solid production-grade infrastructure and backend engineering for a GenAI service, though it lacks complex agentic features.
- Candida Hyphae Segmentation: depth 6.0/10 — The project shows significant depth in model architecture design and data pipeline optimization for medical imaging.
- NetoAI Solutions - Voice Agent: depth 8.0/10 — This project involves complex real-time orchestration and low-latency engineering for a production-grade voice agent.
- NetoAI Solutions - Sales Copilot: depth 10.0/10 — This is a highly sophisticated agentic system featuring multi-agent orchestration, complex tool use, RAG, and rigorous testing.

#2 PRAJWAL A S — 88.2/100

Project: AI Voice Appointment Assistant: A real-time, multilingual voice AI agent for healthcare appointment management. It utilizes WebRTC for low-latency interaction and integrates with a PostgreSQL database via FastAPI to ensure

GitHub: last public activity ~47d ago; relevant repos: voice-ai-backend, python-low-level-design, medical-ai-assistant

- AI Voice Appointment Assistant: depth 7.0/10 — The project demonstrates strong backend integration and tool-calling architecture to ensure data consistency in a real-time voice application.
- AI-Powered HPC Monitoring Assistant: depth 8.0/10 — This project features sophisticated multi-agent orchestration and complex data handling for high-performance computing environments.
- CSIR-4PI HPC Infrastructure & AI Systems: depth 10.0/10 — This is a highly mature, production-deployed system featuring rigorous evaluation, complex data pipelines, and multi-stage agentic workflows.
- best-depth 28.0 + breadth 6.0 (3 solid project(s)) + signals 5.2

#3 V Sree Raghu Vardhan — 88.0/100

Project: AI-driven enterprise HRMS workflows: A complex system of multi-step AI agents integrated with enterprise APIs, semantic retrieval, and document processing to automate HR tasks. It features structured output validation, R

GitHub: No GitHub profile on resume

- AI-driven enterprise HRMS workflows: depth 10.0/10 — This is a highly sophisticated production-grade system featuring multi-step orchestration, state management, rigorous evaluation pipelines, and complex backend integra
- Business Requirements Document generation platform: depth 5.0/10 — The project demonstrates solid backend engineering and integration of LLMs for a specific business use case, though it lacks the advanced agentic complex
- best-depth 28.0 + breadth 4.0 (2 solid project(s)) + signals 6.0

#4 Hari Shanker Sharma — 86.5/100

Project: Sasta Signals: A distributed price monitoring system that tracks products across 7 major Indian e-commerce platforms using automated web scraping. It features real-time price drop detection with notifications and a robust

GitHub: No GitHub profile on resume

- 500Agents: depth 9.0/10 — Demonstrates high-level agentic orchestration using LangGraph, multi-provider routing, and session persistence at scale.
- Jug.ai: depth 8.0/10 — Includes complex features like content training pipelines, real-time intent detection, and AI-to-human handoff logic.
- Support.cc: depth 7.0/10 — A functional business-logic-heavy application integrating RAG-like knowledge management with automated triage workflows.
- TheAIBull.com: depth 9.0/10 — High-stakes system integrating real-time financial data, complex backtesting engines, and secure multi-broker API orchestration.

#5 ARUNIMA SAHA — 85.6/100

Project: Multi-Agent Travel Planning System: An agentic travel planning and booking system that uses LangGraph to orchestrate specialized nodes for flight/hotel search and booking. It features human-in-the-loop workflows, persist

GitHub: No GitHub profile on resume

- Multi-Agent Travel Planning System: depth 9.0/10 — This is a sophisticated multi-agent system featuring complex orchestration, human-in-the-loop workflows, and robust state management.
- Graph RAG Conversational Assistant: depth 8.0/10 — The project demonstrates advanced RAG techniques using knowledge graphs and a multi-stage pipeline for query generation and result formatting.
- AI-Based Text-to-SQL Chatbot: depth 8.0/10 — This project shows strong engineering maturity through its decoupled serverless architecture, security practices, and robust logging/validation layers.
- best-depth 25.2 + breadth 6.0 (3 solid project(s)) + signals 4.5

#6 ABHINAV MISHRA — 85.0/100

Project: Construct IDE: An AI-native teaching IDE that generates personalized coding lessons, implementation tasks, and concept checks. It maintains persistent learner state and session continuity across project steps using LangG

GitHub: last public activity ~0d ago; 60 event(s) in 90d; relevant repos: hackathon-datahub-schemaguard, fxpnt, xitecoin

- Construct IDE: depth 8.0/10 — The project demonstrates sophisticated agentic orchestration and state management to provide a continuous, personalized learning experience.
- Decipath AI: depth 7.0/10 — This project combines web-based retrieval and agentic grounding with complex frontend visualization logic.
- best-depth 22.4 + breadth 4.0 (2 solid project(s)) + signals 3.8

#7 Gourab Das — 85.0/100

Project: SMILE: A full-stack platform built with Django and React to manage donor and patient data. It includes automated infrastructure deployment using Docker and CI/CD pipelines.

GitHub: last public activity ~2d ago; 1 event(s) in 90d

- Tool-augmented AI agents: depth 8.0/10 — This project utilizes advanced orchestration frameworks and tool-calling to solve real-world intelligence gathering tasks.
- Document-processing and Semantic Search Pipeline: depth 7.0/10 — The project shows depth in data engineering and retrieval optimization, specifically addressing chunking and reranking.

- best-depth 22.4 + breadth 4.0 (2 solid project(s)) + signals 3.8

#8 Vivek Chimnani — 83.3/100

Project: AI Research Copilot (LangGraph): A research assistant utilizing a multi-node workflow for planning, research, and report generation. It features session persistence, conditional routing, and modular backend APIs.

GitHub: GitHub rate limit already hit this run

- AI Research Copilot (LangGraph): depth 8.0/10 — The project demonstrates advanced agentic design through multi-node orchestration, conditional routing, and state management.
- best-depth 22.4 + breadth 2.0 (1 solid project(s)) + signals 4.5

#9 Prathamesh Patil — 82.2/100

Project: Adaptive Agentic RAG - Multi-Agent Research System: A multi-agent research system using LangGraph for conditional routing between vector stores and web search. It features self-correction nodes for hallucination detection

GitHub: No GitHub profile on resume

- Adaptive Agentic RAG - Multi-Agent Research System: depth 9.0/10 — The project demonstrates high-level engineering through multi-agent orchestration, self-correction loops, persistent state management, and production-grade
- AI-SaaS-Kit - Scalable Multimodal AI Boilerplate: depth 7.0/10 — This project shows strong backend engineering and infrastructure skills, though it lacks the specific agentic complexity found in the candidate's other work
- best-depth 25.2 + breadth 4.0 (2 solid project(s)) + signals 6.0
- penalty: tutorial/boilerplate: AI-SaaS-Kit - Scalable Multimodal AI Boilerplate (-3.0)

#10 Sumaiya Sultana Shaik — 81.7/100

Project: Privacy Risk Intelligence Platform: A 3-layer privacy-risk analysis system that combines NLP for PII detection, a Random Forest classifier for risk scoring, and a LangChain-orchestrated LLM layer for remediation advice.

GitHub: No GitHub profile on resume

- Privacy Risk Intelligence Platform: depth 8.0/10 — The project demonstrates a sophisticated multi-layer architecture with selective LLM invocation, rigorous testing, and a published methodology.
- WealthWise: AI Financial Planning Platform: depth 7.0/10 — This project shows strong integration of RAG and agentic workflows with a full-stack implementation and custom data engineering.
- best-depth 22.4 + breadth 4.0 (2 solid project(s)) + signals 5.2

Rejected / unparsed (9)

Candidate	Reason
AABHAR GUPTA	No evidence of Python stack; No AI/agentic project evidence
ANJALI PATIL	No evidence of Python stack; No AI/agentic project evidence
Gayathri Palraj	No AI/agentic project evidence
Harshit Kumar Mishra	No evidence of Python stack; No AI/agentic project evidence
Kartikay Sinha	No evidence of Python stack; No AI/agentic project evidence
MD AZIZUL ISLAM	No evidence of Python stack
Mukul Pal	No evidence of Python stack
SHIVAM RAJ	No evidence of Python stack; No AI/agentic project evidence
Shreya Chauhan	No AI/agentic project evidence

Eligibility is a hard, deterministic filter (Python AND AI/agentic evidence). Scores are evidence-source weighted; the LLM contributes AI project depth with cited evidence. GitHub rate-limited/errored profiles are tagged and given a neutral placeholder.